

AMS 2016: UniCloud, Docker at Unidata

LDM, TDS, and RAMADDA on Microsoft Azure VM

Julien Chastang (UCAR, Unidata)

<2015-12-15 Tue>

Contents

1	Preamble	2
2	Quick Start	3
3	Preliminary Setup on Azure	3
3.1	Install <code>docker-machine</code>	4
3.2	Create a VM on Azure.	4
3.3	Configure Unix Shell to Interact with New Azure VM.	4
3.4	Restart Azure VM	4
3.5	<code>ssh</code> into VM with <code>docker-machine</code>	5
3.6	Install Package(s) with <code>apt-get</code>	5
3.7	Add <code>ubuntu</code> User to <code>docker</code> Group and Restart Docker	5
3.8	Install <code>docker-compose</code> on VM	5
4	LDM and TDS Configuration	6
4.1	Background	6
4.1.1	Unidata-Dockerfiles	6
4.1.2	TDSConfig	6
4.2	<code>git clone</code> Repositories	6
4.3	Configuring the LDM	7
4.3.1	LDM Directories on Docker Host	7
4.3.2	LDM Configuration Files	7
4.3.3	Upstream Data Feed from Unidata or Elsewhere	9
4.4	Configuring the TDS	10
4.4.1	Edit TDS catalog.xml Files	10

5	Setting up Data Volumes	10
5.1	Check Free Disk Space	11
5.2	Create /data Directory	11
6	Opening Ports	11
7	Tomcat Logging for TDS and RAMADDA	12
8	Starting the LDM TDS RAMADDA TDM	12
8.0.1	RAMADDA Preconfiguration	12
8.0.2	Final Edit to <code>docker-compose.yml</code>	13
8.0.3	Pull Down Images from the DockerHub Registry . . .	13
8.0.4	Start the LDM, TDS, TDM, RAMADDA	13
9	Check What is Running	13
9.1	Docker Process Status	14
9.2	TDS and RAMADDA URLs	14
9.3	Viewing Data with the IDV	14
9.3.1	Access TDS with the IDV	14
9.3.2	Access RAMADDA with the IDV	14
10	Appendix	15
10.1	Common Problems	15
10.1.1	Certificate Regeneration	15
10.1.2	Size of Image is not Large Enough	15

1 Preamble

This guide describes how to configure the LDM¹, TDS², and RAMADDA³ on a Microsoft Azure VM⁴. The document assumes you have access to Azure resources though these instructions should be fairly similar on other cloud providers (e.g., Amazon). They also require familiarity with Unix, Docker, and Unidata technology in general. You must have `sudo` privileges on the Azure host which will hopefully be available you. You must be comfortable entering commands at the Unix command line. We will be using Docker images defined at the Unidata-Dockerfiles repository⁵ in addition to a con-

¹<http://www.unidata.ucar.edu/software/ldm/>

²<http://www.unidata.ucar.edu/software/thredds/current/tds/>

³<http://sourceforge.net/projects/ramadda/>

⁴<https://azure.microsoft.com>

⁵<https://github.com/Unidata/Unidata-Dockerfiles>

figuration specifically planned for an AMS 2016 demonstrations project⁶.

2 Quick Start

In order to best understand this configuration process, it is recommended to read the complete contents of this document and follow the instructions starting in the next section. If there are problems you will be able to reason about the errors. However, if you are eager to get started, you can follow this quick start section.

- `git clone https://github.com/Unidata/Unidata-Dockerfiles`
- Download and install⁷ `docker-machine`
- Run the `Unidata-Dockerfiles/ams2016/unicloud-1.sh` script (this will take few minutes) to create the Docker host on Azure. See the section on Azure for more information.

For example,

```
unicloud-1.sh --azure-host <azure-host> --azure-subscription-id "3.14" \  
--azure-subscription-cert "/path/to/mycert.pem"
```

Now you are ready to do additional configuration on the new Docker host:

```
docker-machine ssh <azure-host> "bash -s" < \  
Unidata-Dockerfiles/ams2016/unicloud-2.sh
```

`ssh` into new Docker host with `docker-machine ssh <azure-host>`. See the section below about editing the `ldmfile.sh` to correctly handle logging. Run `~/git/Unidata-Dockerfiles/ams2016/unicloud-3.sh`. See the section on checking what you have done.

3 Preliminary Setup on Azure

The VM we are about to create will be our **Docker Host** from where we will run Docker containers for the LDM, TDS, and RAMADDA.

⁶<https://github.com/Unidata/Unidata-Dockerfiles/tree/master/ams2016>

⁷<https://docs.docker.com/machine/install-machine/>

3.1 Install docker-machine

Install⁸ `docker-machine` on your local computer. `docker-machine` is a command line tool that gives users the ability to create Docker VMs on your local computer or on a cloud provider such as Azure.

3.2 Create a VM on Azure.

The following `docker-machine` command will create a Docker VM on Azure for running various Unidata Docker containers. **Replace the environment variables with your choices.** This command will take a few minutes to run (between 5 and 10 minutes). You will have to supply `azure-subscription-id` and `azure-subscription-cert` path. See the Azure `docker-machine` instructions⁹, if you have questions about this process. Also set the size of the VM¹⁰ (e.g., `Small`, `ExtraLarge`) and supply the name of the Azure Docker host.

```
# Create Azure VM via docker-machine
docker-machine -D create -d azure \
    --azure-subscription-id=$AZURE_ID \
    --azure-subscription-cert=$AZURE_CERT \
    --azure-size=$AZURE_SIZE $AZURE_HOST
```

3.3 Configure Unix Shell to Interact with New Azure VM.

Execute the following `eval` command on your local computer shell environment to ensure that `docker` commands will be run with the newly created Docker host.

```
# Ensure docker commands will be run with new host
eval "$(docker-machine env $AZURE_HOST)"
```

3.4 Restart Azure VM

Mysteriously, when you `ssh` (see next section) into the fresh VM, you are immediately told to restart it so let's preempt that message by doing that now.

⁸<https://docs.docker.com/machine/install-machine/>

⁹<https://azure.microsoft.com/en-us/documentation/articles/virtual-machines-docker-machine/>

¹⁰<https://azure.microsoft.com/en-us/documentation/articles/virtual-machines-size-specs/>

```
# immediately restart VM, according to Azure
docker-machine restart $AZURE_HOST
# Again, ensure docker commands will be run with new host
eval "$(docker-machine env $AZURE_HOST)"
```

3.5 ssh into VM with docker-machine

```
docker-machine ssh $AZURE_HOST
```

3.6 Install Package(s) with apt-get

At the very least, we will need `unzip` on the Azure Docker host.

```
# update and install package(s)
sudo apt-get -qq update
sudo apt-get -qq install unzip
```

3.7 Add ubuntu User to docker Group and Restart Docker

```
# Add ubuntu to docker group
sudo usermod -G docker ubuntu
```

```
# Restart docker service
sudo service docker restart
```

In Unix, when adding a user to a group, it is simply easiest to logout and log back in for this change to be recognized. Do that by exiting the VM and logging back in with `docker-machine ssh` command.

3.8 Install docker-compose on VM

`docker-compose` is a tool for defining and running multi-container Docker applications. In our case, we will be running the LDM, TDS, TDM (THREDDDS Data Manager) and RAMADDA so `docker-compose` is perfect for this scenario. Install `docker-compose` on the Azure Docker host. You may have to update version (currently at 1.5.2).

```
# Get docker-compose
curl -L \
https://github.com/docker/compose/releases/download/1.5.2/docker-compose-$(uname -s)-$(uname -m) > docker-compose
sudo mv docker-compose /usr/local/bin/
sudo chmod +x /usr/local/bin/docker-compose
```

4 LDM and TDS Configuration

4.1 Background

We have done the preliminary legwork to tackle the next step in this process. We will now want to clone two repositories that will allow us to configure and start running the the LDM, TDS, and RAMADDA. In particular, we will be cloning:

- github.com/Unidata/Unidata-Dockerfiles¹¹
- github.com/Unidata/TdsConfig¹²

4.1.1 Unidata-Dockerfiles

The `Unidata-Dockerfiles` repository contains a number of Dockerfiles that pertain to various Unidata technologies (e.g., the LDM) and also projects (e.g., `ams2016`). As a matter of background information, a `Dockerfile` is a text file that contains commands to build a Docker image containing, for example, a working LDM. These Docker images can subsequently be run by `docker` command line tools, or `docker-compose` commands that rely on a `docker-compose.yml` configuration file. A `docker-compose.yml` file is a text file that captures exactly how one or more containers run including directory mappings (from outside to within the container), port mappings (from outside to within the container), and other information.

4.1.2 TDSConfig

The `TDSConfig` repository is a project that captures THREDDs and LDM configuration files (e.g., `catalog.xml`, `pqact.conf`) for the TDS at <http://thredds.ucar.edu>. Specifically, these TDS and LDM configurations were meant to work in harmony with one another. We can re-use this configuration with some minor adjustments for running the TDS on the Azure cloud.

4.2 git clone Repositories

With that background information out of the way, let's clone those repositories by creating `~/git` directory where our repositories will live and issuing some `git` commands.

¹¹<https://github.com/Unidata/Unidata-Dockerfiles>

¹²<https://github.com/Unidata/TdsConfig>

```
# Get the git repositories we will want to work with
mkdir -p /home/ubuntu/git
git clone https://github.com/Unidata/Unidata-Dockerfiles \
    /home/ubuntu/git/Unidata-Dockerfiles
git clone https://github.com/Unidata/TdsConfig /home/ubuntu/git/TdsConfig
```

4.3 Configuring the LDM

4.3.1 LDM Directories on Docker Host

For anyone who has worked with the LDM, you may be familiar with the following directories:

- etc/
- var/data
- var/logs
- var/queue

The LDM **etc** directory is where you will find configuration files related to the LDM including `ldmd.conf`, `pqact` files, `registry.xml`, and `scour.conf`. We will need the ability to easily observe and manipulate the files from **outside** the running LDM container. To that end, we need to find a home for **etc** on the Docker host. The same is true for the **var/data** and **var/logs** directories. Later, we will use Docker commands that have been written on your behalf to mount these directories from **outside** to **within** the container. The **var/queues** directory will remain inside the container.

```
# Create LDM directories
mkdir -p ~/var/logs
mkdir -p ~/etc/TDS
```

var/data is a bit different in that it needs to be mounted on data volume on the Docker host. We will be handling that step further on.

4.3.2 LDM Configuration Files

There is a generic set of LDM configuration files located here `~/git/Unidata-Dockerfiles/ldm/etc/`. However, we will just grab `netcheck.conf` which will remain unmodified.

```
# Copy various files for the LDM.
cp ~/git/Unidata-Dockerfiles/ldm/etc/netcheck.conf ~/etc
```

The rest of the LDM configuration files will come from our `ams2016` project directory.

Also, remember that these files will be used **inside** the LDM container that we will set up shortly. We will now be working with these files:

- `ldmd.conf`
- `registry.xml`
- `scour.conf`

1. `ldmd.conf`

```
cp ~/git/Unidata-Dockerfiles/ams2016/ldmd.conf ~/etc/
```

This `ldmd.conf` has been setup for the AMS 2016 demonstration serving the following data feeds:

- 13km Rapid Refresh¹³
- NESDIS GOES Satellite Data¹⁴
- Unidata NEXRAD Composites

For your information, and for future reference, there is a `~/git/TdConfig/idd/pqacts/README.txt` file that may be helpful in writing a suitable `ldmd.conf` file.

2. `registry.xml`

```
cp ~/git/Unidata-Dockerfiles/ams2016/registry.xml ~/etc/
```

This file has been set up for the AMS 2016 demonstration. Otherwise you would have to edit the `registry.xml` to ensure the `hostname` element is correct. For your own cloud VMs, work with `support-idd@unidata.ucar.edu` to devise a correct `hostname` element so that LDM statistics get properly reported. Here is an example `hostname` element: `unidata-server.azure.unidata.ucar.edu`.

3. `scour.conf`

You need to scour data or else your disk will full up. The crontab entry that runs scour is in the LDM Docker container¹⁵. Scouring is invoked once per day.

¹³<http://rapidrefresh.noaa.gov/>

¹⁴http://www.nesdis.noaa.gov/imagery_data.html

¹⁵<https://github.com/Unidata/Unidata-Dockerfiles/blob/master/ldm/crontab>


```
cp ~/git/Unidata-Dockerfiles/ams2016/scour.conf ~/etc/
```

4. pqact.conf and TDS configuration

In the `ldmd.conf` file we copied just a moment ago there is a reference to a `pqact` file; `etc/TDS/pqact.forecastModels`. We need to ensure that file exists by doing the following instructions. Specifically, explode `~/git/TdsConfig/idd/config.zip` into `~/tdsconfig` and `cp -r` the `pqacts` directory into `~/etc/TDS`. **Note** do NOT use soft links. Docker does not like them.

```
# Set up LDM and TDS configuration
mkdir -p ~/tdsconfig/
cp ~/git/TdsConfig/idd/config.zip ~/tdsconfig/
unzip ~/tdsconfig/config.zip -d ~/tdsconfig/
cp -r ~/tdsconfig/pqacts/* ~/etc/TDS
```

5. Edit `ldmfile.sh`

Open the `etc/TDS/util/ldmfile.sh` file in the editor of your choice. As the top of this file indicates, you must edit the `logfile` to suit your needs. Change the

```
logfile=logs/ldm-mcidas.log
```

line to

```
logfile=var/logs/ldm-mcidas.log
```

This will ensure `ldmfile.sh` can properly be invoked from the `pqact` files.

4.3.3 Upstream Data Feed from Unidata or Elsewhere

The LDM operates on a push data model. You will have to find someone who will agree to push you the data. If you are part of the American academic community please send a support email to `support-idd@unidata.ucar.edu` to discuss your LDM data requirements.

4.4 Configuring the TDS

4.4.1 Edit TDS catalog.xml Files

The `catalog.xml` files for TDS configuration are contained within the `~/tdsconfig` directory. Search for all files terminating in `.xml` in that directory. Edit the `xml` files for what data you wish to server. See the TDS Documentation¹⁶ for more information on editing these XML files.

Let's see what is available in the `~/tdsconfig` directory.

```
find ~/tdsconfig -type f -name "*.xml"

/home/ubuntu/tdsconfig/idd/forecastModels.xml
/home/ubuntu/tdsconfig/idd/radars.xml
/home/ubuntu/tdsconfig/idd/obsData.xml
/home/ubuntu/tdsconfig/idd/forecastProdsAndAna.xml
/home/ubuntu/tdsconfig/idd/satellite.xml
/home/ubuntu/tdsconfig/radar/CS039_L2_stations.xml
/home/ubuntu/tdsconfig/radar/CS039_stations.xml
/home/ubuntu/tdsconfig/radar/RadarNexradStations.xml
/home/ubuntu/tdsconfig/radar/RadarTerminalStations.xml
/home/ubuntu/tdsconfig/radar/RadarL2Stations.xml
/home/ubuntu/tdsconfig/radar/radarCollections.xml
/home/ubuntu/tdsconfig/catalog.xml
/home/ubuntu/tdsconfig/threddsConfig.xml
/home/ubuntu/tdsconfig/wmsConfig.xml
```

5 Setting up Data Volumes

As alluded to earlier, we will have to set up data volumes so that the LDM can write data, and the TDS and RAMADDA can have access to that data. The `/mnt` has lots of space, but the storage is considered **ephemeral** so be careful! Azure makes no effort to backup data on `/mnt`. For the LDM this should not be too much of a problem because real time data is coming in and getting scoured continuously, but for for any other application you may wish to be careful as there is the potential to lose data. There is more information about this topic here¹⁷.

¹⁶<http://www.unidata.ucar.edu/software/thredds/current/tds/catalog/index.html>

¹⁷<https://azure.microsoft.com/en-us/documentation/articles/virtual-machines-linux-how-to-attach-disk/>

5.1 Check Free Disk Space

Let's first display the free disk space with the `df` command.

```
df -H
```

Filesystem	Size	Used	Avail	Use%	Mounted	on
/dev/sda1	31G	2.0G	28G	7%	/	
none	4.1k	0	4.1k	0%	/sys/fs/cgroup	
udev	7.4G	8.2k	7.4G	1%	/dev	
tmpfs	1.5G	394k	1.5G	1%	/run	
none	5.3M	0	5.3M	0%	/run/lock	
none	7.4G	0	7.4G	0%	/run/shm	
none	105M	0	105M	0%	/run/user	
none	66k	0	66k	0%	/etc/network/interfaces.dynamic.d	
/dev/sdb1	640G	73M	607G	1%	/mnt	

5.2 Create /data Directory

Create a `/data` directory where the LDM can write data soft link to the `/mnt` directory. Also, create a `/repository` directory where RAMADDA data will reside.

```
# Set up data directories
sudo ln -s /mnt /data
sudo mkdir /mnt/ldm/
sudo chown -R ubuntu:docker /data/ldm
sudo mkdir /home/ubuntu/repository/
sudo chown -R ubuntu:docker /home/ubuntu/repository
```

These directories will be used by the LDM, TDS, and RAMADDA docker containers when we mount directories from the Docker host into these containers.

6 Opening Ports

Ensure these ports are open¹⁸ on the VM where these containers will run.

¹⁸<https://azure.microsoft.com/en-us/documentation/articles/virtual-machines-set-up-endpoints/>

Service	External Port
HTTP	80
TDS	8080
RAMADDA	8081
SSL TDM	8443
LDM	388

Note the TDM is an application that works in conjunction with the TDS. It creates indexes for GRIB data in the background, and notifies the TDS via port 8443 when data have been updated or changed. See here¹⁹ to learn more about the TDM.

7 Tomcat Logging for TDS and RAMADDA

It is a good idea to mount Tomcat logging directories outside the container so that they can be managed for both the TDS and RAMADDA.

```
# Create Tomcat logging directories
mkdir -p ~/logs/ramadda-tomcat
mkdir -p ~/logs/tds-tomcat
```

Note there is also a logging directory in `~/tdsconfig/logs`. All these logging directories should be looked at periodically, not the least to ensure that log files are not filling up your system.

8 Starting the LDM TDS RAMADDA TDM

8.0.1 RAMADDA Preconfiguration

When you start RAMADDA for the very first time, you must have a `password.properties` file in the RAMADDA home directory which is `/home/ubuntu/repository/`. See RAMADDA documentation²⁰ for more details on setting up RAMADDA. Here is a `pw.properties` file to get you going. Change password below to something more secure!

```
# Create RAMADDA default password
echo ramadda.install.password=changeme! > /home/ubuntu/repository/pw.properties
```

¹⁹<https://www.unidata.ucar.edu/software/thredds/current/tds/reference/collections/TDM.html>

²⁰<http://ramadda.org/repository/userguide/toc.html>

8.0.2 Final Edit to docker-compose.yml

When the TDM communicates to the TDS concerning changes in data it observes with data supplied by the LDM, it will communicate via the `tdm` tomcat user. Edit the `docker-compose.yml` file and change the `TDM_PW` to `MeIndexer`. This is not as insecure as it would seem since the `tdm` user has few privileges. Optimally, one could change the password hash for the TDM user in the `tomcat-users.xml` file.

8.0.3 Pull Down Images from the DockerHub Registry

You are almost ready to run the whole kit and caboodle. But first pull the relevant docker images to make this easier for the subsequent `docker-compose` command.

```
# Docker pull all relevant images
docker pull unidata/ldmtds:latest
docker pull unidata/tdm:latest
docker pull unidata/tds:latest
docker pull unidata/ramadda:latest
```

8.0.4 Start the LDM, TDS, TDM, RAMADDA

We are now finally ready to start the LDM, TDS, TDM, RAMADDA with the following `docker-compose` command.

```
# Start up all images
docker-compose -f ~/git/Unidata-Dockerfiles/ams2016/docker-compose.yml up -d
```

9 Check What is Running

In this section, we will assume you have created a VM called `unidata-server`. You should have these services running:

- LDM
- TDS
- TDM
- RAMADDA

Next, we will check our work through various means.

9.1 Docker Process Status

From the shell where you started `docker-machine` earlier you can execute the following `docker ps` command to list the containers on your docker host. It should look something like the output below.

```
docker ps --format "table {{.ID}}\t{{.Image}}\t{{.Status}}"
```

CONTAINER	ID	IMAGE	STATUS
4ed1c4c18814	unidata/ramadda:latest	Up	17 seconds
bdfcf5590bc6	unidata/ldmtds:latest	Up	18 seconds
aee044cf8e66	unidata/tdm:latest	Up	20 seconds
4d0208f85b22	unidata/tds:latest	Up	21 seconds

9.2 TDS and RAMADDA URLs

Verify what you have the TDS and RAMADDA running by, for example, navigating to: <http://unidata-server.cloudapp.net/thredds/catalog.html> and <http://unidata-server.cloudapp.net:8081/repository>. If you are going to RAMADDA for the first time, you will have to do some RAMADDA set up²¹.

9.3 Viewing Data with the IDV

Another way to verify your work is run the Unidata Integrated Data Viewer²².

9.3.1 Access TDS with the IDV

In the IDV Dashboard²³, you should be able to enter the catalog XML URL: <http://unidata-server.cloudapp.net/thredds/catalog.xml>.

9.3.2 Access RAMADDA with the IDV

RAMADDA has good integration with the IDV and the two technologies work well together. You may wish to install the RAMADDA IDV plugin²⁴ to publish IDV bundles to RAMADDA. RAMADDA also has access to the `/data/ldm` directory so you may want to set up server-side

²¹<http://ramadda.org/repository/userguide/toc.html>

²²<https://www.unidata.ucar.edu/software/idv/>

²³<https://www.unidata.ucar.edu/software/idv/docs/userguide/data/choosers/CatalogChooser.html>

²⁴<http://www.unidata.ucar.edu/software/idv/docs/workshop/savingstate/>

view of this part of the file system²⁵. Finally, you can enter this catalog URL in the IDV dashboard to examine data holdings shared bundles, etc. on RAMADDA `http://unidata-server.cloudapp.net:8081/repository?output=thredds.catalog`.

10 Appendix

10.1 Common Problems

10.1.1 Certificate Regeneration

When using `docker-machine` may see an error message pertaining to regenerating certificates.

```
Error running connection boilerplate: Error checking and/or regenerating the certs: The
You can attempt to regenerate them using 'docker-machine regenerate-certs name'.
Be advised that this will trigger a Docker daemon restart which will stop running conta
```

In this case:

```
docker-machine regenerate-certs <azure-host>
eval "$(docker-machine env <azure-host>)"
```

Like the error message says, you may need to restart your Docker containers with `docker-compose`, for example.

10.1.2 Size of Image is not Large Enough

If you see your containers not starting or error messages like this:

```
ERROR: Cannot start container ef229d1753b24b484687ac4d6b8a9f3b961f2981057c59266c45b9d5
```

it is possible you did not create a sufficiently large VM. Try increasing the size of the VM²⁶.

Ramadda.html

²⁵<http://ramadda.org//repository/userguide/developer/filesystem.html>

²⁶<https://azure.microsoft.com/en-us/documentation/articles/virtual-machines-size-specs/>